

Operating Systems

Recitation 6

Plan

- Multitasking
 - Single process
 - Cooperative multitasking
 - Preemptive multitasking
- Processes API (Linux)

What is a process?

- Process == task
- Execution of an assignment defined by a program
 - Process is instance of program execution
 - Two opened notepads == two processes
 - But not straightforward - several Firefox windows may be supported by single process

What is a process?


- For OS
 - “Context of execution”
 - independent entity consuming resources (CPU time, memory, disk)
- How does the OS switch between processes?

Single-process system

- Running program gets full control over computer resources
- Next program starts after previous ended
- What's the problem?

Single-process system

- Not efficient
 - CPU idle a lot of the time – for example when waiting for printer to finish
 - May want to prioritize



Cooperative Multitasking

Programs voluntarily give up CPU time for other programs

- Program gets control only when other program “yields”
- User may select/activate program to run
 - Pause editor and look into spreadsheet
- Examples:
 - Windows 3.1, old Mac OS
- Today not in common use. Why?

Cooperative Multitasking

Difficult for programmers to write a program aware of other programs

- How does a program know when to yield CPU to other program, and when (if ever) will it resume?
- We may need access to private memory of other processes
 - To know when and how to yield...
 - Lots of IPC

Preemptive multitasking

OS Scheduler decides when to run each process and for how long

- CPU time divided to slices
- Process gets fair share of CPU time (priority & policy)
- OS manages each process state, guards its private memory and manages shared resources (disk, printer,...)

Preemptive multitasking

- Process may enter “waiting state” (e.g. wait for other process/event to complete) so OS won’t allocate any CPU time at all
- Process suspended without its permission (i.e. preemption)
- OS manages queues of prioritized processes

Multitasking: Some Requirements

- To switch between programs, OS able to
 - Interrupt process on intervals
 - Freeze state of paused program (e.g. current instruction, register and memory state)
 - And resume later
 - Use “scheduler” to decide what to do next according to some policy


Process descriptor

Each process has a “descriptor” (defined in `struct task_struct` in `/usr/src/linux-headers-XXX/include/linux/sched.h`):

- Process Identifier (pid)
- Program counter - Address in memory of next instruction to execute
- Context data - Registers
- Statistics - CPU time used till now
- State - Running, ready, waiting, ...
- Priority - Low, high
- Open files
- Memory map

Scheduling

- In class saw example of well-known scheduler implementation – Multilevel feedback queues
- A Unix OS manages several queues of process descriptors
 - Doubly linked list, Priorities



מעבר בין תורים

❖ מעבר בין תורים:

- תהליך שצבר זמן ריצה רב עובר לתור נמוך יותר
- תהליך שלא רץ זמן רב עובר לתור גבוה יותר, למניעת הרעבה
- תהליך שחיכה למשאב וקיבל אותו נכנס לתור גבוה מאוד כדי שישחרר את המשאב מהר ככל האפשר

❖ העדיפות שאנו מעניקים לתהליכים משפיעה על הפרמטרים של מעבר בין תורים ולא קובעת תור הספציפי לתהליך


❖ קל לממש מדיניות הוגנת ויעילה

❖ קל לממש יכולות מיוחדות, כגון תהליכים שרצים רק שאף תהליך אחר לא מוכן לריצה

Multitasking API

- Process created by `fork()` from “parent” process
- Linux starts with 2 processes
 - Swapper (manages memory)
 - `init` (`pid=1`) – all processes are its sons (forked)
- “Parent” can fork multiple times
 - Son too...

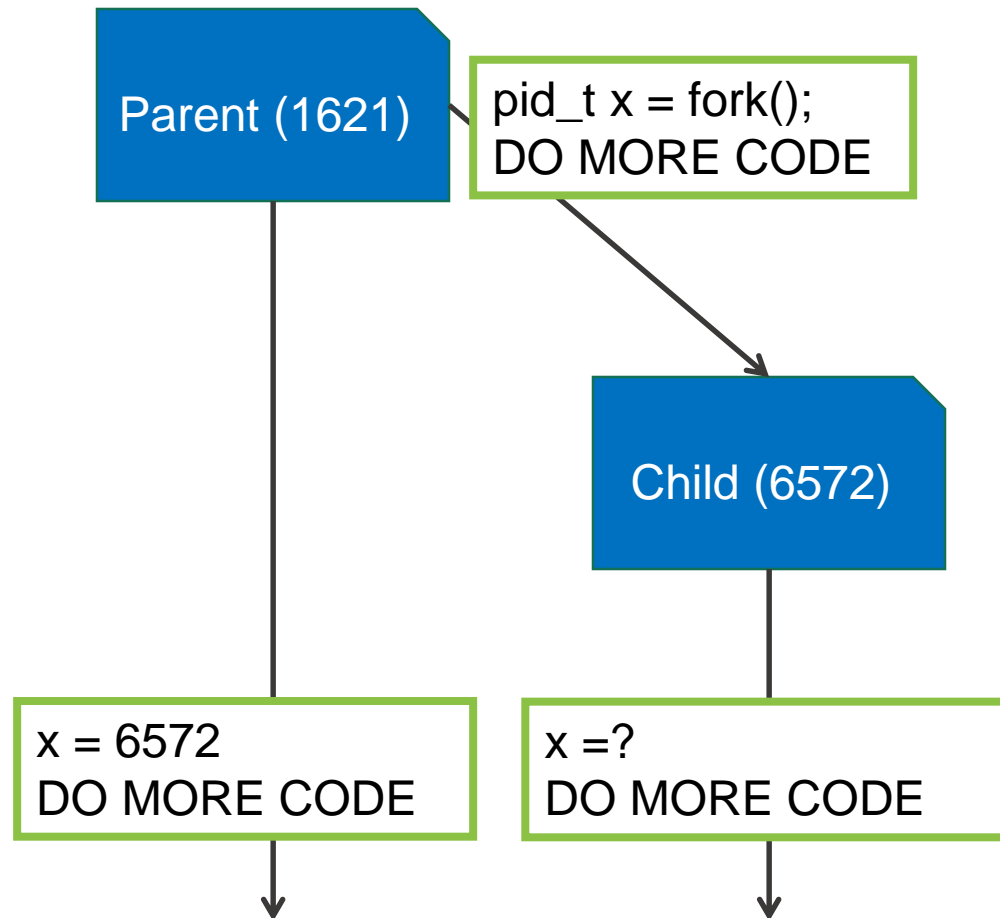
fork()



`pid_t fork()` - returns identifier of process

- Copies parent to son (clone)
 - Code and execution point
 - Environment = open files and fd's, current working lib
 - Memory = buffers, variables and values
- Return value
 - Son = 0, parent = son's new pid, error = -1
- Windows? – `CreateProcess` (sort of `fork+execv`)

fork()



fork()

- Typical fork usage example

```
int pid = fork();  
if (pid < 0)  
    // fork() failed - handle error  
  
if (pid == 0)  
    // son process - do son code  
else  
    // parent process - do parent code
```

Peculiarities (1)

- What would this do?

```
main() {  
    int pid1 = fork();  
    int pid2 = fork();  
}
```

- How many processes?

Peculiarities (2)

- What would this output?

```
main() {  
    fork();  
    fork();  
    printf("hello");  
}
```

- “hellohellohello”?

Peculiarities (3)

- What would this do?

```
main() {  
    int pid = fork();  
    if (pid) {  
        fork();  
    }  
    fork();  
}
```

- How many processes?

execv()

```
int execv(const char *filename, char
*const argv[])
```

- Run a program in same process
 - filename - file with program to run
 - argv - same as in main (argv[0] == filename)
 - Last entry in argv should be NULL
 - On failure -1, **on success doesn't return!**

execv()

- Typical usage example

```
main() {  
    char *argv[] = {"/bin/date", "-u", NULL};  
    execv("/bin/date", argv);  
    printf("hello");  
}
```

- What would this do?
- ***Code example***

`wait()` / `waitpid()`

```
pid_t wait(int *status)
```

```
pid_t waitpid(pid_t pid, int *status,  
int options)
```

- Parent can wait for son to end (any son / specific)
 - `status` - son's status when finished
 - Returns -1 if all sons finished or already `wait()`ed, or **`pid`** of some process that finished

wait() / waitpid()

- How do we wait for all child processes to end?

```
while (wait(&status) != -1);
```

- `pid_t getpid();`
 - Get process id of current process

- ***Code example***

Zombies

- Sometimes need information from son-process, or can't continue until it finishes → need to wait for it
- ...if son is finished, and `wait()` not invoked yet, it is a **“Zombie”**.
 - Exists only as process record
 - Deleted once son `wait()`ed
- What if parent is finished, and son isn't?
- ***Code example***



Linux Signals

- Inform process async. on some event
- Form of IPC and for messages from OS about relevant events
- 32 types in UNIX
- Some signals are handled automatically by OS
 - SIGKILL, SIGSTOP
- Others handled by signal handler of process
 - OS queues signals (only 1 for every type!)
 - Upon returning to user-mode process handles them using handler
 - Can write our own signal handlers too!

kill()

```
int kill(pid_t pid, int sig);
```

- Officially - signal process to die
 - `pid` = process id
 - `sig` = signal to send
- Actually name is misleading
 - can send any signal
 - Real kill signal intercepted by OS...
- **kill** utility from command line really does send a SIGKILL / SIGTERM signal