

Operating Systems

Recitation 9

Plan

- Networking background
 - Protocols
 - Layers
 - TCP/IP stack - Ethernet, TCP, DNS etc.
- Sockets & API (Linux)
- Client/Server scheme


What is the Internet

Internet: “network of networks”

- loosely hierarchical
- public Internet Vs private intranet

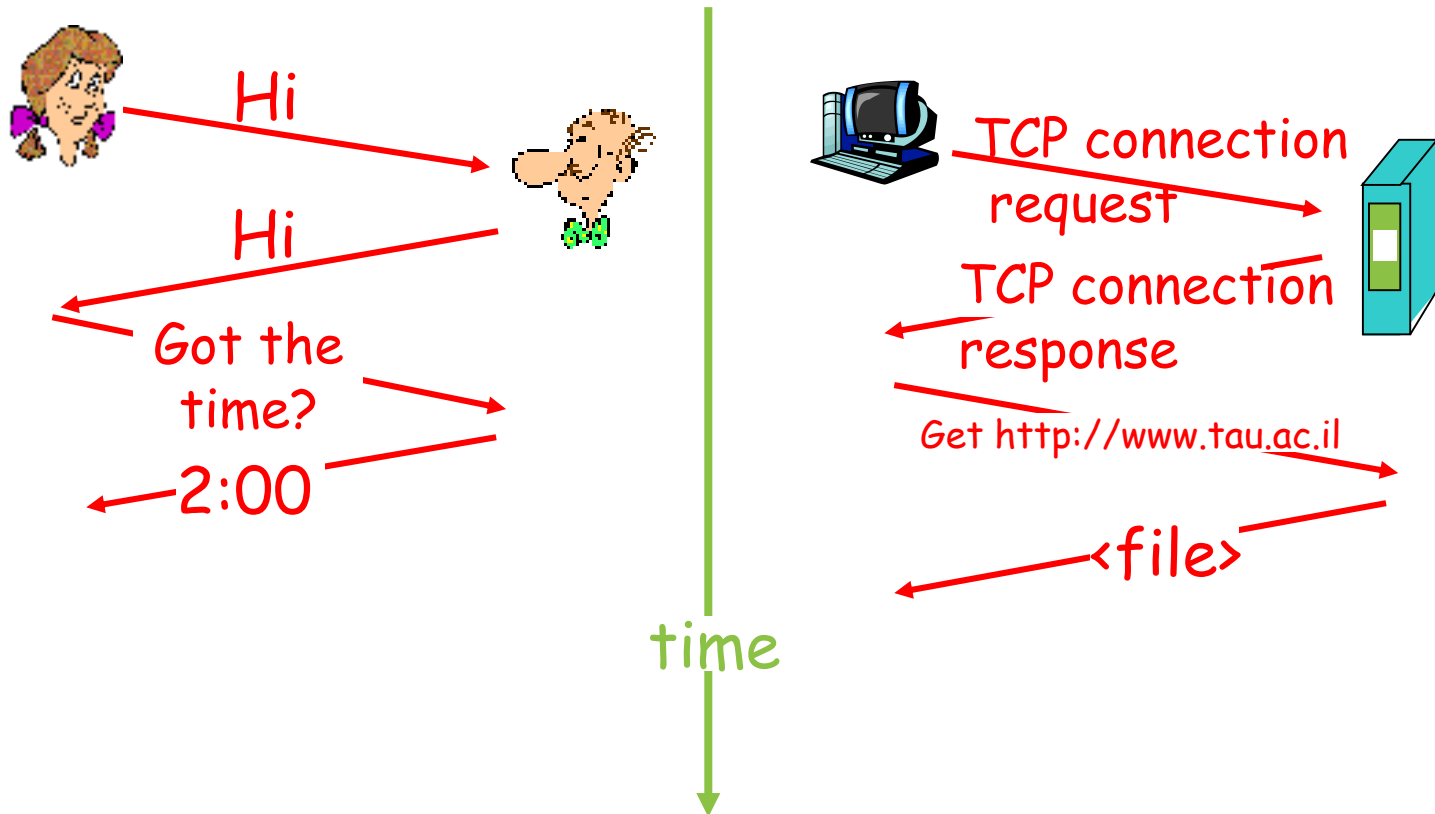
protocols coordinate communication

- Who gets to transmit?
- What path to take?
- What message format?
- e.g., TCP, IP, HTTP, FTP



Protocols

- A human protocol and a computer network protocol:

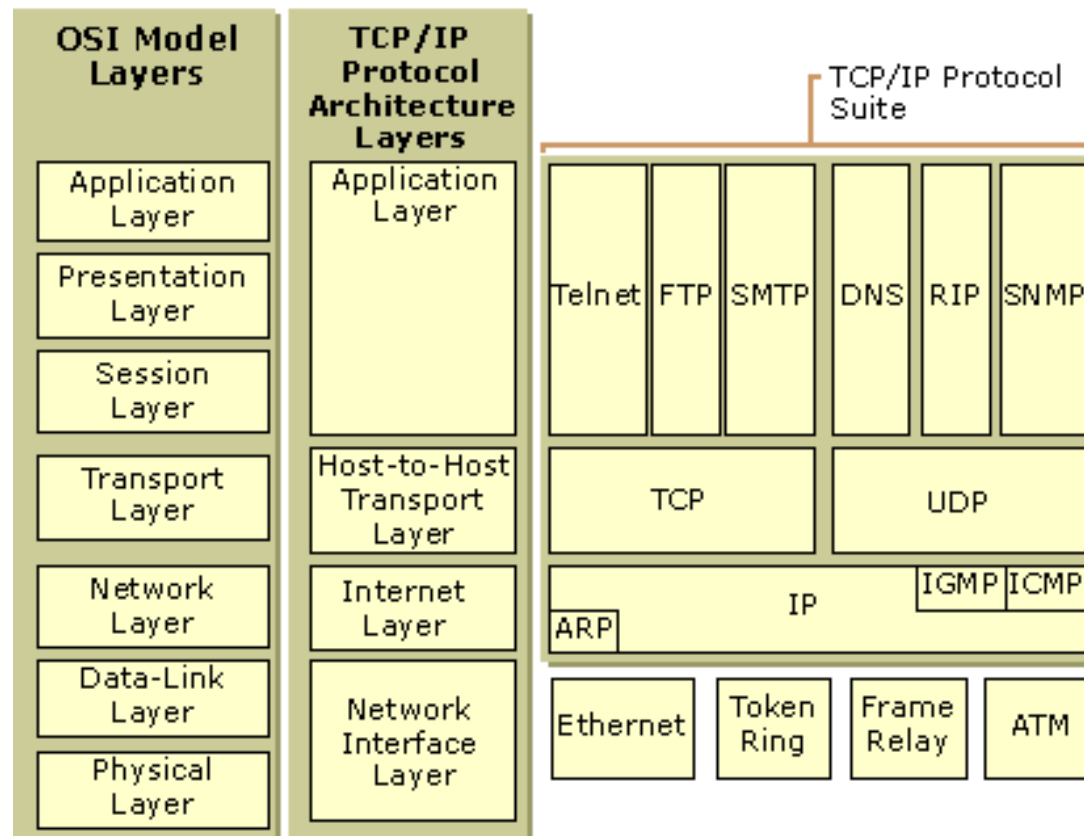


Layering

- Protocols are “stack”-ed in **Layers:**
- Each layer implements a service
- Same layers on each side of connection communicate data to each other
 - Layers rely on services provided by layer below

TCP/IP protocol stack

- Most common protocol stack




Ethernet: Bus


- Network interface layer (lowest)
- Several Stations connected by wire (bus), (“everyone sees everyone”)
- Each network card station has unique address
 - MAC == Medium Access Control
- Need to know each other’s address on the physical line to communicate data to it
- Nothing is guaranteed



From Computer Desktop Encyclopedia
© 1999 The Computer Language Co., Inc.



A network card, network adapter, LAN Adapter or NIC



IP address

- Internet layer
- Numerical address (115.64.32.12-human readable form)
- Global unique addresses to be reached from outside world
- Divided to sub-networks



Subnets

- IP address:
 - subnet part (high order bits)
 - host part (low order bits)
- *What's a subnet ?*
 - device interfaces with same subnet part of IP address
 - can physically reach each other without intervening router
- Local addresses to use within local networks can be reused (10.0.2.X, 192.168.1.X...)



200.23.16.0/23

IP Routing

- Packets reach their destination **Hop-by-hop**
- Think of how you send a package through the mail from Tel Aviv to Beer Sheva
 - local mail office (TA)
 - regional branch (TA)
 - regional branch (BS)
 - local mail office (BS)
- How do we know where to deliver next (next “hop”)?:
 - Use Routing tables

IP is “Best-Offer”

- No guarantees
- Lots of things can go wrong:
 - data corruption
 - lost data packets
 - duplicate arrival
 - out-of-order packet delivery
- To improve, IP packet “data” portion needs to implement some protocol too
→ i.e. transport layer (TCP, UDP)



TCP over IP (TCP/IP)

- **Transmission Control Protocol (TCP)**

Offsets Octet		0								1								2								3								
Octet	Bit	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	
0	0	Source port																Destination port																
4	32	Sequence number																																
8	64	Acknowledgment number (if ACK set)																																
12	96	Data offset	Reserved 0 0 0			N S		C W R E	E C R G	U R C K	A C S H	P S S T	R S Y N	S Y N	F I N	Window Size																		
16	128	Checksum																Urgent pointer (if URG set)																
20	160	Options (if Data Offset > 5, padded at the end with "0" bytes if necessary)																																
...																																

TCP over IP (TCP/IP)

- Guaranteed delivery
 - Packets have sequence number (SQN)
 - Receiver send acknowledge (ACK)
- Port number in addition to IP address
 - To distinguish between applications
- Mechanisms for establishment and termination of a connection

The Importance of Ports

- Ports == 16 bit identifiers
- Uniquely identify processes involved in a socket (TCP & UDP)
- Help route data to destination
- In UNIX first 1024 ports for both protocols are called “well known ports”
- Defined in /etc/services
- Programs that bind to these ports require “root” access

UDP over IP (User Datagram)

Offsets	Octet	0								1								2								3							
Octet	Bit	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
0	0	Source port																Destination port															
4	32	Length																Checksum															

- Simple transport-layer protocol
- Over IP
- Ports to distinguish applications at the same host
- Length & checksum for verification

UDP over IP (User Datagram)

- No guaranteed delivery
- May arrive out of order (different routing path)
- Light Weight (less delay, processing)
- Usage?
 - Audio streaming
 - Why?

Sockets

- Yet another form of IPC
- Linux/Windows use sockets to communicate between processes over the network
 - remote, but also local!
- Socket is a concept
 - represented in software by its handle
- In Unix its a “file”
 - usually nameless
 - represents a network connection == source/dest IPs & ports + protocol
- Implemented in socket library
 - `#include <sys/types.h>`
 - `#include <sys/socket.h>`

socket()

```
int socket(int family, int type, int protocol)
```

- “creates an endpoint for communication and returns a descriptor” - man 2 socket
- family - protocol family
 - AF_INET (IPv4 Internet protocols)
 - AF_UNIX (local socket = local IPC!)
- type - communication semantics
 - SOCK_STREAM - full duplex, reliable connection
 - SOCK_DGRAM - datagrams, unreliable, connectionless

socket() cont.

```
int socket(int family, int type, int protocol)
```

- protocol – transport layer protocol to use. 0 chooses default
 - For type=SOCK_STREAM, protocol 0 is TCP
 - What about SOCK_DGRAM?
- Protocols for AF_UNIX? none 😊
- Returns - socket (file) descriptor

Binding a socket

```
int bind(int sockfd, struct sockaddr  
*my_addr, int addrlen)
```

- sockfd – returned from socket()
- my_addr – address to bind to, machine IP and desired port
 - Example soon
- addrlen – my_addr length in bytes
- Returns - 0 on success, 1 failure

Listening to a socket

```
int listen(int sockfd, int num)
```

- Mark socket as a “special” - one used to accept incoming connection requests
 - SOCK_DGRAM not legal. Why?
- sockfd – the socket
- num – max # of pending connections
- Returns – 0 success, 1 failure

Accepting a Connection

```
int accept(int sockfd, const struct sockaddr*  
          serv_addr, socklen_t addrlen)
```

- Wait on socket for incoming requests.
 - SOCK_DGRAM not legal...
- Create **new socket** with new descriptor
- sockfd – listening socket
- serv_addr – *remote* server address
 - Example soon
- addrlen - # of bytes in serv_addr
- Returns – new sockfd on success, -1 failure

Connecting to a Socket

```
int connect(int sockfd, const struct sockaddr*  
            serv_addr, socklen_t addrlen)
```

- sockfd – socket
- serv_addr – *remote* server address
 - Example soon
- addrlen - # of bytes in serv_addr
- Returns – 0 success, 1 failure

Address data structures

- Format and size of my_addr is usually protocol specific.
- struct sockaddr is used as the “base” of a set of address structures

```
struct sockaddr {  
    unsigned short sa_family; // address family, AF_xxx  
    char sa_data[14]; // 14 bytes of protocol address  
};
```


Address data structures (2)

```
struct sockaddr_in {
    short int sin_family; // address family, AF_xxx
    unsigned short int sin_port; // port number
    struct in_addr sin_addr; // internet address
    unsigned char sin_zero[8]; // for alignments
};

struct in_addr {
    unsigned long s_addr; //32-bit long (4 bytes) IP
                           address
}
```

Typical Usage

```
struct sockaddr_in serverName;  
serverName.sin_family = AF_INET;  
...  
connect(clientSocket,  
        (struct sockaddr*) &serverName,  
        sizeof(serverName));
```

Communicating data

```
int send(int sockfd, const void *msg,  
         size_t len, int flags)  
int recv(int sockfd, void *buf,  
         size_t len, int flags)
```

- Send data, with special options and flags.
- Conveniently can also be done like regular files with `read()/write()` on `sockfd`
- By default blocking = control returned to process only when write/read occurs
 - Though its configurable

read()/write() semantics

- read()
 - blocks until data is available
 - may return $<$ than max bytes (whatever is available)
 - you must be prepared to read data 1 byte at a time!
 - **How do we know when to finish?**
- write()
 - might not be able to write all bytes
 - returns # of bytes sent (not received...)

Close socket

```
int close(int sockfd)
```

- Close socket descriptor pointed by sockfd
- Either end of the connection can close()
- If other end closed connection, and there is no buffered data, reading from a TCP socket returns 0 to indicate EOF.
- Returns – 0 success, 1 failure


client	server
sd=socket()	sd=socket()
	bind(sd, port)
	listen(sd,...)
connect(sd, dst)	new_sd=accept(sd)
write(sd, ...)	write(new_sd, ...)
read(sd, ...)	read(new_sd, ...)
close(sd)	close(new_sd)

Some explaining regarding client/server

- Why do we need `bind()`?
 - Otherwise server listens on random port. Clients don't know where to connect!
- Why do we need `listen`?
 - Possibly multiple clients send requests
 - `listen` defines length of queue for such incoming connections
- Main thread does all the work here!
 - May want to fork or delegate work to threads, so main thread can continue servicing incoming connections

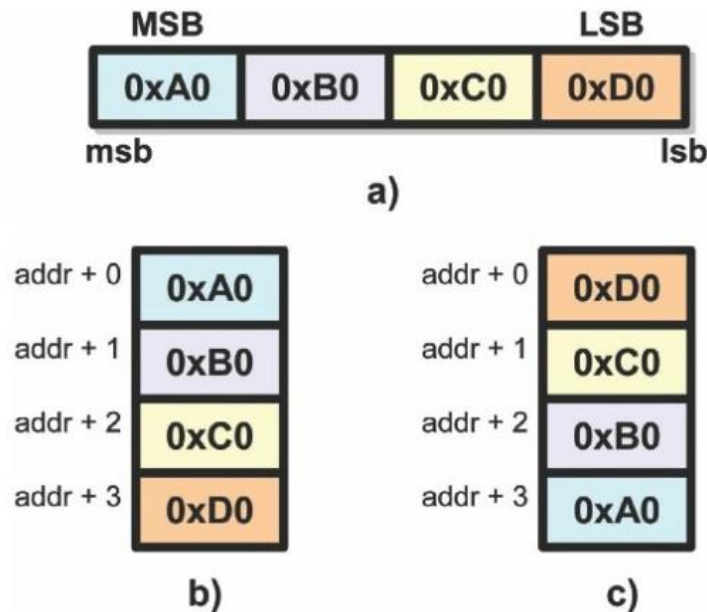
Endianness

- One last note before code
- A long long time ago in Lilliput, society split into two factions!
 - Big-Endians - opened soft-boiled eggs at larger end ("the primitive way")
 - Little-Endians - broke their eggs at the smaller end.
- The feud continues...



Endianness

- Different CPUs have different order of bytes – MSB, LSB
 - **Big-Endian** - MSB of any multibyte data field is stored at the lowest memory address which is also the address of the larger field.
 - **Little-Endian** – same but for LSB



Endianness

- Why do we care? we communicate between remote computers, with possibly different CPUs
 - X86 little endian
- When we pass numbers, need uniform representation

- Socket library utilities

```
u_long htonl(u_long); // host to network long (32 bits)  
u_short htons(u_short); // host to network short (16 bits)
```

```
u_long ntohl(u_long); // network to host long (32 bits)  
u_short ntohs(u_short); // network to host short (16 bits)
```

Example

- TCP server/client
- ***Code example***